

Maximilian Knespel  
Center for Information Services and High Performance Computing (ZIH)

# Ratarmount – A Tool for Mounting Large TAR Files

2020, June

# Why?

Why bother with mounting a TAR instead of extracting it?

- The file system might not support millions of very small files very well
- There might not be sufficient space for extracting
- The extracted uncompressed data might be multiple times larger
- You only need read-only access to a few files for previewing

# Why?

Why create a new tool?

Archivemount <https://github.com/cybernoid/archivemount/commits/master> exists for this but:

- Mounting can take hours
- No progress indicator during mounting
- Each subsequent mount will again take hours
- Accessing files can take more than a second
- Memory usage as long as it is mounted can be tens of gigabytes
- No integrated extended features like: recursive mounting, union mounting, support for updated files in the TAR, ...

# Ratarmount - RAndom Access TAR Mount

<https://github.com/mxmlnkn/ratarmount>

Main functions as required for mounting  
ImageNet:

- FUSE layer for easy access to files in TAR
- Find the offsets for each file to seek to
- Write out offsets and metadata to a file for reuse on subsequent mounting
- Keep memory usage low by streaming directly to and from the index file
- Support mounting TARs inside TARs recursively

```
fall11_whole.tar
├─ n00004475.tar
├─ n00005787.tar
├─ n00006024.tar
├─ n00006484.tar
└─ ...
n00005787.tar
├─ n00005787_13.JPEG
├─ n00005787_32.JPEG
├─ n00005787_54.JPEG
├─ n00005787_58.JPEG
└─ ...
```

↓ **ratarmount**

```
mounted_fall11_whole
├─ n00004475.tar
│   └─ n00004475_6590.JPEG
│   └─ ...
├─ n00005787.tar
│   └─ n00005787_13.JPEG
│   └─ ...
└─ ...
```

fall11\_whole.index.sqlite

# Ratarmount

## Additional features

Started out as a fairly simple ~300 lines python script.

These are some additional things initiated by reported issues on GitHub

- Support more obscure TAR features like sparse files and hard links
- Travis CI and automatic uploads to PyPI
- Support for bz2 and gzip compressed TAR files
- Mount not only tar.bz2 but also simple bz2 compressed files for uncompressed access
- Avoid memory leaks in tarfile module and by using SQLite for the index
- Work with incomplete TAR files
- Access older overwritten versions of a file analogous to tar --occurrence=<n> ...
- Union mounting and bind mounting

Now ~1700 lines, plus tests, plus the external indexed\_bzip2 repository for bzip2 decompression.


# Ratarmount

## Union Mounting and file versions


Specify more than one input to get a merged view. An input folder **can** at the same time also be an output folder.

Access old versions of a <file> using the special hidden <file>.versions subfolder. The highest version number is the latest and the one shown by default.

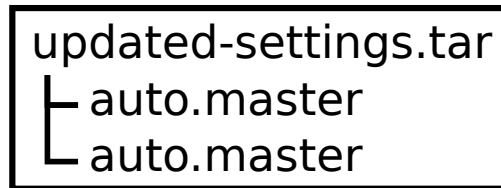
First input: compressed TAR



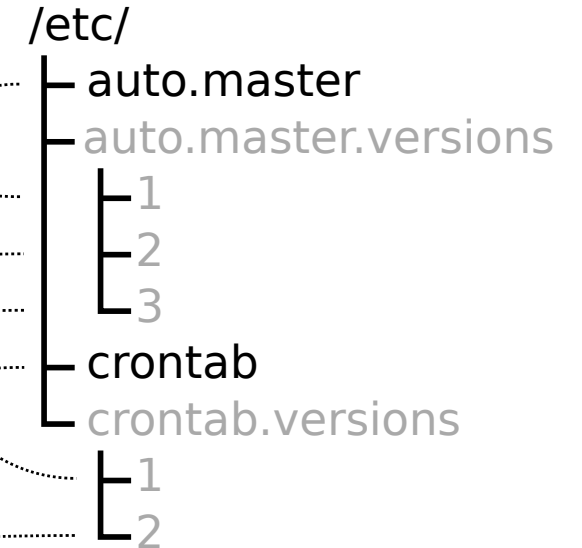

**ratarmount**



Second input: simple TAR



Third input: folder




# How does it work?

The tape archive format (TAR) is pretty simple. It basically contains the all files concatenated to each other aligned to 512B and interspersed with headers.

In order to seek to an arbitrary file, collect metadata and offsets:

File Path	Seek Offset	File Size
n00004475.tar/ n00004475_6590.JPEG	1024	78483
n00004475.tar/ n00004475_15899.JPEG	80384	94671

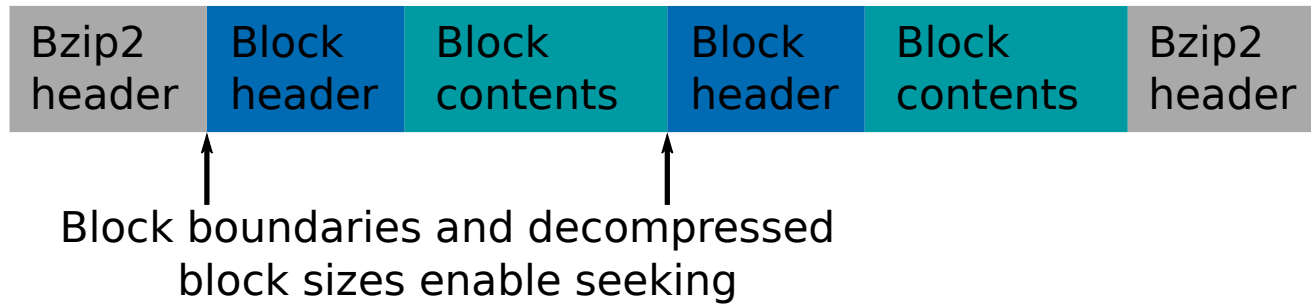
This is the bare minimum but other metadata like user name, id, ... are also stored for FUSE.



Field Offset	Field Size	Field
0	156	(as in old format)
156	1	Type flag
157	100	(as in old format)
257	6	USTAR indicator "ustar"
263	2	USTAR version "00"
265	32	Owner user name
297	32	Owner group name
329	8	Device major number
337	8	Device minor number
345	155	Filename prefix

# How does it work?

## Bzip2 Seeking



Collect the block header offsets with bit precision and the corresponding offsets in the decoded data. The latter is the accumulated sum of all decompressed block sizes before it.

I modified the 0BSD licensed Bzip2 decoder from <https://github.com/landley/toybox> by Robert Landley to collect, export and import these offsets and added a Python interface in [https://github.com/mxmInkn/indexed\\_bzip2](https://github.com/mxmInkn/indexed_bzip2)

Gzip seeking done similarly with [https://github.com/pauldmccarthy/indexed\\_gzip](https://github.com/pauldmccarthy/indexed_gzip) by Paul McCarthy


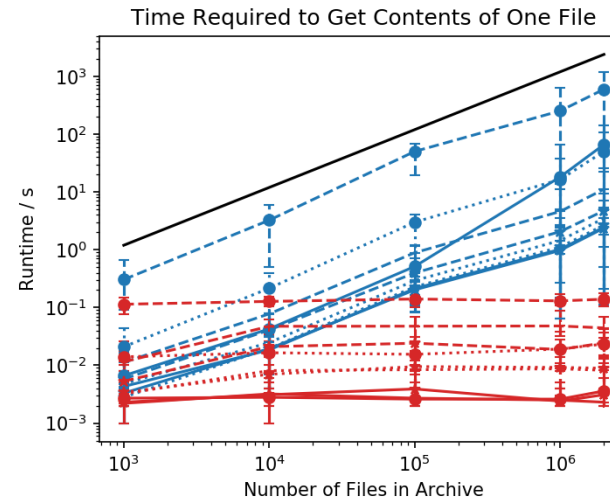
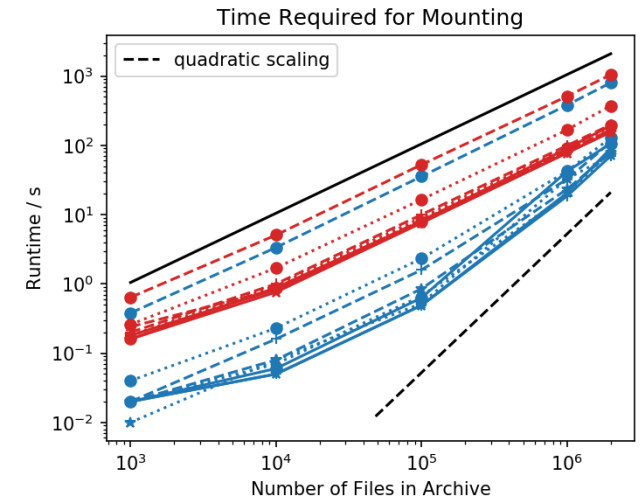



# Benchmarks

## Comparison with archivemount

Storing and loading indexes alone is already a huge advantage. But, here are some other comparisons:

- Ratarmount memory usage is bounded except for the gzip backend.
- Mounting time unexpectedly is slower than archivemount at least for < 2M files. Scaling suggests an inversion.
- Accessing a file scales with the number of files when using archivemount!
- Tree traversal with the find command is for some reason 10x slower in ratarmount



# Benchmarks

## Index Serialization Backends

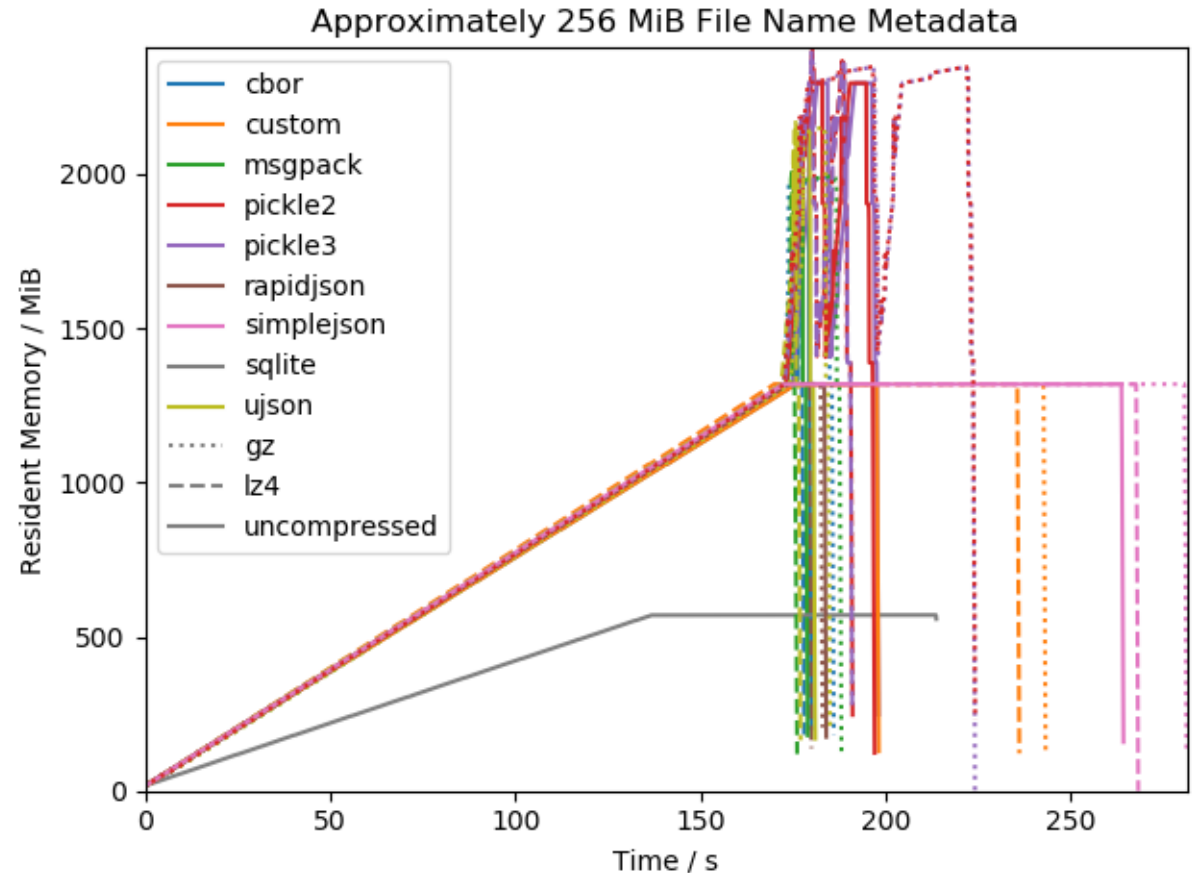
Periodically measure the resident memory size for different backends, which all save the same metadata.

Two phases:

- Metadata aggregation in memory
- Metadata dumping

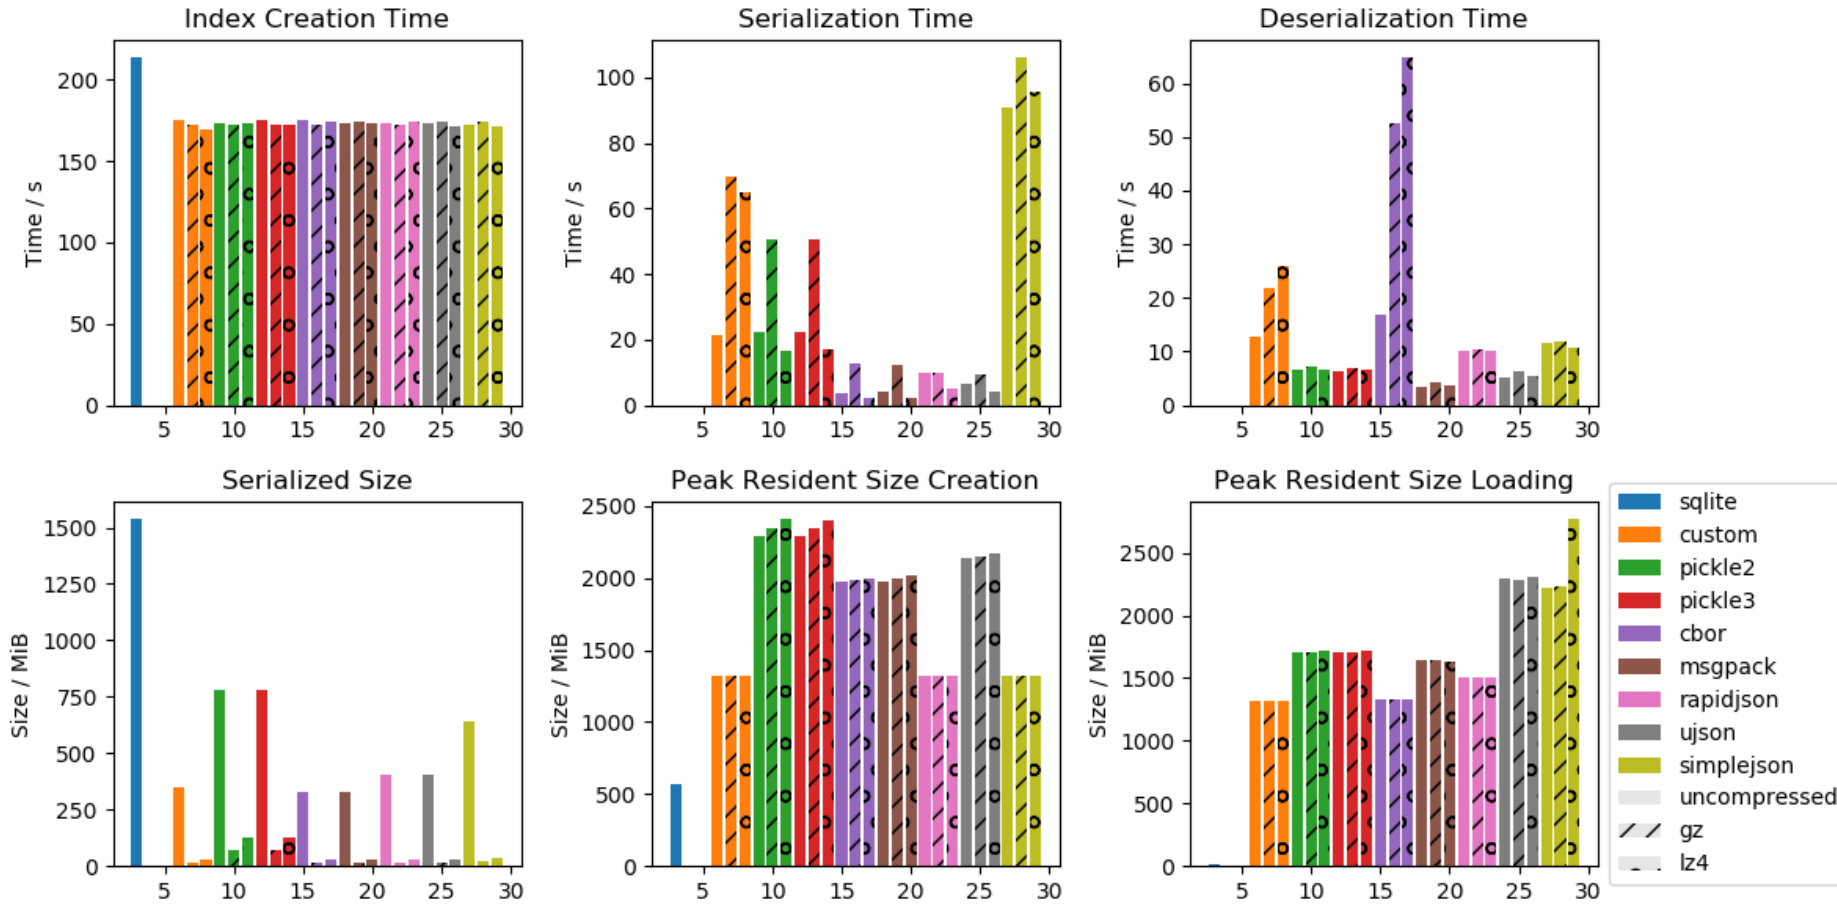
Results:

- Only SQLite has bounded memory usage
- Many backends double their memory usage during actually writing data



# Benchmarks

## Index Serialization Backends



# Future work

## Ideas and ToDos:

- Optimize: sequential reading, bzip2 decoding, tree traversal with find
- Parallelize the Bzip2 decoder using one thread per bzip2 block and some kind of prediction
- Parallelize TAR indexing?
- Option for a writable overlay layer similar to Singularity has
- Limited write support, which would append updated and new files to the TAR?
- Add xz support? <https://github.com/tomorrow-nf/tar-as-filesystem> by Kyle Davidson and Tyler Morrow seems to be able to do this using a modified xz decoder.
- Support other systems than Linux, e.g., by using reFUSE a fork of fusepy